

```
In [74]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [75]: df=pd.read_csv(r"C:\Users\balakumar\OneDrive\Desktop\dataset\work.csv")
```

```
In [76]: df
```

```
Out[76]:
```

	Res. No	Age	Gender	Marital status	Education	Income	Covid- infected or not	CAS1	CAS2	CAS3	CAS4	CAS5
0	1	2	1	1	1	1	1	1	1	2	2	2
1	2	1	0	1	1	2	1	1	1	2	2	2
2	3	2	0	1	3	1	1	3	3	0	0	3
3	4	3	0	0	1	2	1	0	1	0	1	0
4	5	2	0	1	1	1	1	1	2	1	1	2
...
1345	1346	2	0	0	2	2	1	3	2	1	3	3
1346	1347	3	0	0	1	2	1	3	3	3	3	3
1347	1348	3	0	0	1	2	0	1	0	1	1	2
1348	1349	1	1	0	1	2	1	3	2	1	2	3
1349	1350	2	0	1	1	1	0	2	0	0	0	1

1350 rows × 12 columns

```
In [77]: df.shape
```

```
Out[77]: (1350, 12)
```

```
In [78]: df.isnull().sum()
```

```
Out[78]: Res. No          0
Age              0
Gender           0
Marital status   0
Education        0
Income           0
Covid- infected or not 0
CAS1             0
CAS2             0
CAS3             0
CAS4             0
CAS5             0
dtype: int64
```

```
In [79]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1350 entries, 0 to 1349
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Res. No               1350 non-null   int64
 1   Age                   1350 non-null   int64
 2   Gender                1350 non-null   int64
 3   Marital status        1350 non-null   int64
 4   Education              1350 non-null   int64
 5   Income                 1350 non-null   int64
 6   Covid- infected or not 1350 non-null   int64
 7   CAS1                  1350 non-null   object
 8   CAS2                  1350 non-null   int64
 9   CAS3                  1350 non-null   int64
10   CAS4                  1350 non-null   int64
11   CAS5                  1350 non-null   int64
dtypes: int64(11), object(1)
memory usage: 126.7+ KB
```

```
In [80]: df.columns
```

```
Out[80]: Index(['Res. No', 'Age', 'Gender', 'Marital status', 'Education', 'Income',  
              'Covid- infected or not', 'CAS1', 'CAS2', 'CAS3', 'CAS4', 'CAS5'],  
              dtype='object')
```

```
In [81]: #To check correlation between vectors  
correlation = df.corr()
```

```
In [82]: correlation
```

```
Out[82]:
```

	Res. No	Age	Gender	Marital status	Education	Income	Covid- infected or not	CAS2	CAS3	CAS4	CAS5
Res. No	1.000000	0.059598	-0.030123	-0.067225	-0.002179	0.055200	-0.007846	0.015461	0.046159	0.042492	0.014578
Age	0.059598	1.000000	0.044289	-0.080058	-0.058806	0.084014	-0.030567	0.009756	0.042519	0.029378	0.020434
Gender	-0.030123	0.044289	1.000000	0.032329	-0.093072	-0.019621	-0.000768	0.035036	0.022752	0.027767	-0.000631
Marital status	-0.067225	-0.080058	0.032329	1.000000	0.008777	-0.661565	0.043465	0.014007	0.029137	0.026799	-0.054360
Education	-0.002179	-0.058806	-0.093072	0.008777	1.000000	-0.006456	0.021575	-0.002785	-0.013849	-0.012464	0.016278
Income	0.055200	0.084014	-0.019621	-0.661565	-0.006456	1.000000	-0.003209	-0.053953	0.003141	-0.035159	-0.013909
Covid- infected or not	-0.007846	-0.030567	-0.000768	0.043465	0.021575	-0.003209	1.000000	0.020668	-0.012125	-0.022504	-0.044196
CAS2	0.015461	0.009756	0.035036	0.014007	-0.002785	-0.053953	0.020668	1.000000	0.319270	0.370506	0.137341
CAS3	0.046159	0.042519	0.022752	0.029137	-0.013849	0.003141	-0.012125	0.319270	1.000000	0.263840	0.152906
CAS4	0.042492	0.029378	0.027767	0.026799	-0.012464	-0.035159	-0.022504	0.370506	0.263840	1.000000	0.118909
CAS5	0.014578	0.020434	-0.000631	-0.054360	0.016278	-0.013909	-0.044196	0.137341	0.152906	0.118909	1.000000

```
In [83]: gender = pd.crosstab(index = df["Gender"],
                             columns = 'count',
                             normalize = True)

print(gender)
```

col_0	count
Gender	
0	0.561481
1	0.432593
4	0.005185
5	0.000741

```
In [84]: #Covid-19 infected migrant workers with gender proportion
gender_stat = pd.crosstab(index = df['Gender'],columns =df['Covid- infected or not'],margins = True, normalize = 'inde
print(gender_stat)
```

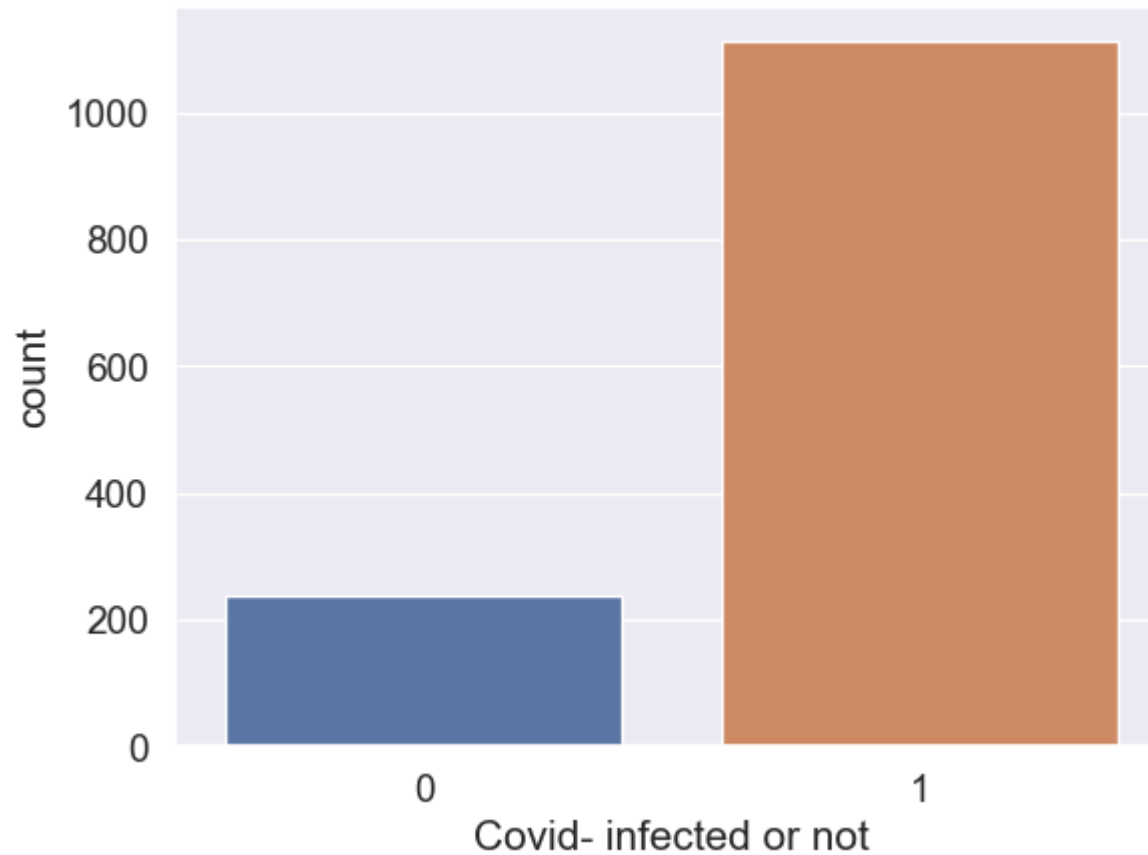
Covid- infected or not	0	1
Gender		
0	0.174142	0.825858
1	0.179795	0.820205
4	0.142857	0.857143
5	0.000000	1.000000
All	0.176296	0.823704

```
In [85]: covid_19_infected = sns.countplot(df['Covid- infected or not'])  
covid_19_infected
```

C:\Users\balakumar\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning:

Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
Out[85]: <AxesSubplot:xlabel='Covid- infected or not', ylabel='count'>
```




```
In [86]: #  
sns.distplot(df['Age'],bins = 5, kde = False)
```

C:\Users\balakumar\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:


`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

Out[86]: <AxesSubplot:xlabel='Age'>




```
In [87]: df["Covid- infected or not"].value_counts().plot(kind="pie")
```

```
Out[87]: <AxesSubplot:ylabel='Covid- infected or not'>
```



```
In [88]: cormat = df.corr()
f, ax = plt.subplots(figsize=(12,9))
sns.heatmap(cormat, vmax=.8, square=True);
plt.show
```

```
Out[88]: <function matplotlib.pyplot.show(close=None, block=None)>
```





2

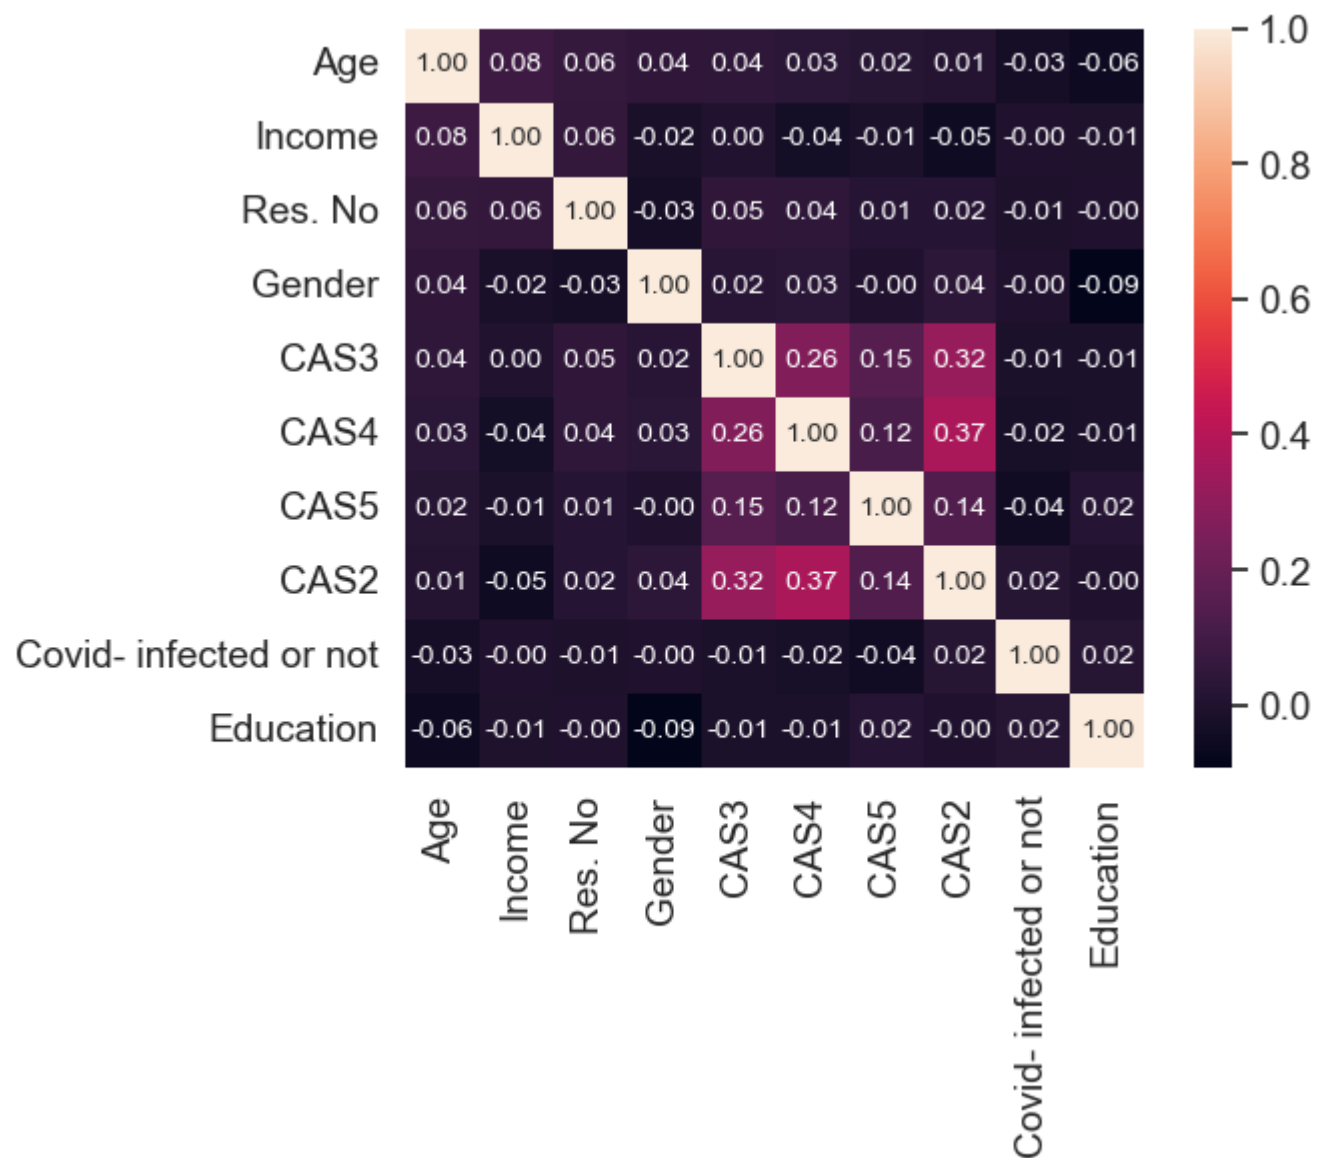
Covid-in

```
In [89]: plt.figure(figsize=(35,20))
sns.heatmap(df.corr(), annot=True, cmap="YlGnBu")
```

Out[89]: <AxesSubplot:>



```
In [90]: #treatment correlation matrix
k = 10
cols = cormat.nlargest(k, 'Age')['Age'].index
cm = np.corrcoef(df[cols].values.T)
sns.set(font_scale=1.25)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_kws={'size': 10}, yticklabels=cols, xticklabels=cols)
plt.show()
```



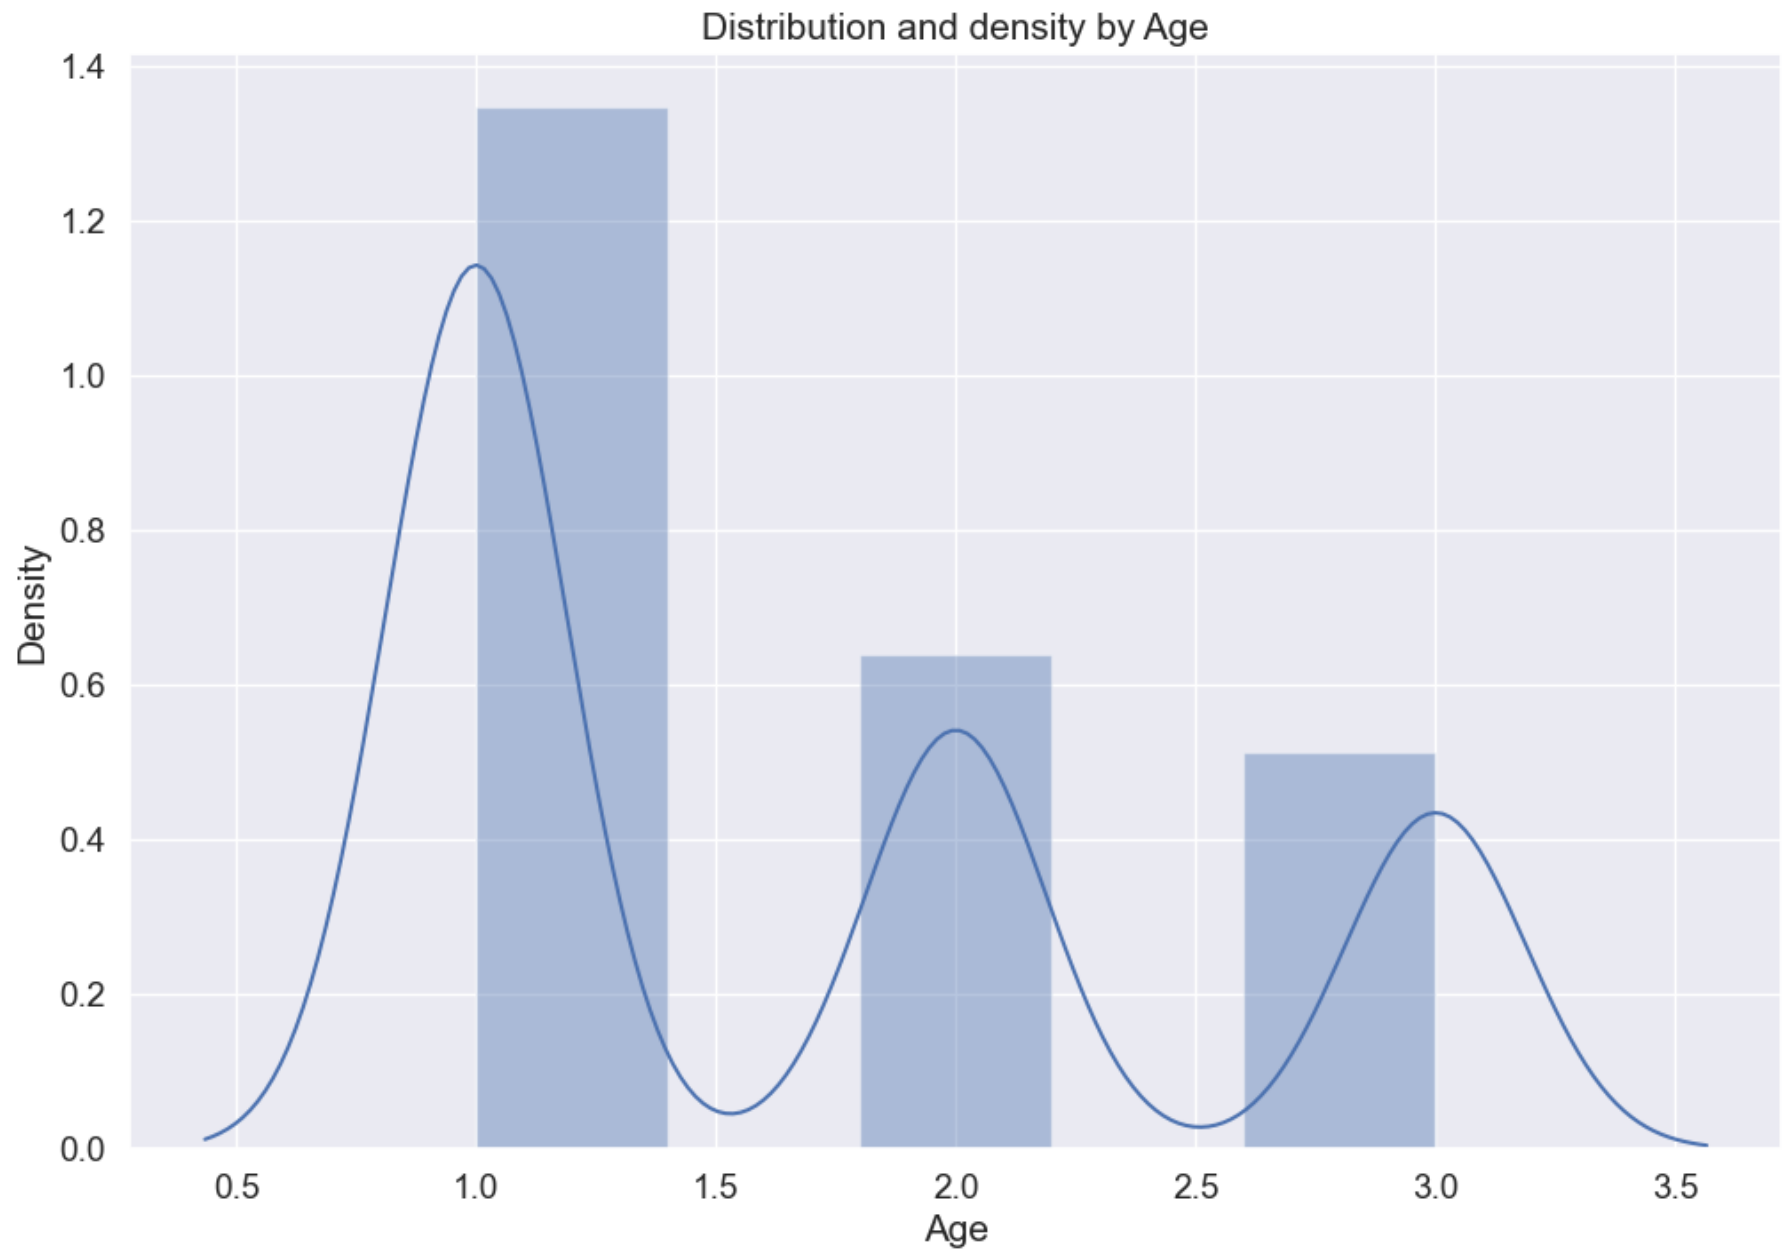
Some charts to see data relationship

```
In [91]: #Distribution and density by Age
plt.figure(figsize=(12,8))
sns.distplot(df['Age'],bins=5)
plt.title("Distribution and density by Age")
plt.xlabel('Age')
```

C:\Users\balakumar\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

Out[91]: Text(0.5, 0, 'Age')




```
In [92]: plt.figure(figsize=(12,8))
labels = df['Age']
g = sns.countplot(x='Covid- infected or not', data = df)
g.set_xticklabels(labels)
plt.title("Total Distribution by Covid impacted or not")
```


```
-----
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_8180\4257121072.py in <module>
      2 labels = df['Age']
      3 g = sns.countplot(x='Covid- infected or not', data = df)
----> 4 g.set_xticklabels(labels)
      5 plt.title("Total Distribution by Covid impacted or not")

~\anaconda3\lib\site-packages\matplotlib\axes\_base.py in wrapper(self, *args, **kwargs)
     73
     74     def wrapper(self, *args, **kwargs):
--> 75         return get_method(self)(*args, **kwargs)
     76
     77     wrapper.__module__ = owner.__module__

~\anaconda3\lib\site-packages\matplotlib\axis.py in _set_ticklabels(self, labels, fontdict, minor, **kwargs)
    1796     if fontdict is not None:
    1797         kwargs.update(fontdict)
-> 1798     return self.set_ticklabels(labels, minor=minor, **kwargs)
    1799
    1800     def _set_tick_locations(self, ticks, *, minor=False):

~\anaconda3\lib\site-packages\matplotlib\axis.py in set_ticklabels(self, ticklabels, minor, **kwargs)
    1718         # remove all tick labels, so only error for > 0 ticklabels
    1719         if len(locator.locs) != len(ticklabels) and len(ticklabels) != 0:
-> 1720             raise ValueError(
    1721                 "The number of FixedLocator locations"
    1722                 f" ({len(locator.locs)}), usually from a call to"
```

ValueError: The number of FixedLocator locations (2), usually from a call to set_ticks, does not match the number of ticklabels (1350).



```
In [93]: o = df["Age"]
g = sns.factorplot(x='Age', y='Covid- infected or not', hue='Gender', data = df, kind="bar", ci=None)
g.set_xticklabels(o)

plt.title('Probability of mental health condition')
plt.ylabel('Probability x 100')
plt.xlabel('Age')

#replace legend labels
new_labels = df["Gender"]
for t,l in zip(g._legend.texts, new_labels):t.set_text(l)

#positioning the legend
g.fig.subplots_adjust(top=0.9, right=0.9)
plt.show()
```

C:\Users\balakumar\anaconda3\lib\site-packages\seaborn\categorical.py:3717: UserWarning:

The `factorplot` function has been renamed to `catplot`. The original name will be removed in a future release. Please update your code. Note that the default `kind` in `factorplot` (`'point'`) has changed `'strip'` in `catplot`.

```

-----
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_8180\791786080.py in <module>
      1 o = df["Age"]
      2 g = sns.factorplot(x='Age', y='Covid- infected or not', hue='Gender', data = df, kind="bar", ci=None)
----> 3 g.set_xticklabels(o)
      4
      5 plt.title('Probability of mental health condition')

~\anaconda3\lib\site-packages\seaborn\axisgrid.py in set_xticklabels(self, labels, step, **kwargs)
    880         ax.set_xticklabels(curr_labels, **kwargs)
    881     else:
--> 882         ax.set_xticklabels(labels, **kwargs)
    883     return self
    884


~\anaconda3\lib\site-packages\matplotlib\axes\_base.py in wrapper(self, *args, **kwargs)
     73
     74     def wrapper(self, *args, **kwargs):
---> 75         return get_method(self)(*args, **kwargs)
     76
     77     wrapper.__module__ = owner.__module__

~\anaconda3\lib\site-packages\matplotlib\axis.py in _set_ticklabels(self, labels, fontdict, minor, **kwargs)
    1796     if fontdict is not None:
    1797         kwargs.update(fontdict)
-> 1798     return self.set_ticklabels(labels, minor=minor, **kwargs)
    1799
    1800     def _set_tick_locations(self, ticks, *, minor=False):

~\anaconda3\lib\site-packages\matplotlib\axis.py in set_ticklabels(self, ticklabels, minor, **kwargs)
    1718         # remove all tick labels, so only error for > 0 ticklabels
    1719         if len(locator.locs) != len(ticklabels) and len(ticklabels) != 0:
-> 1720             raise ValueError(
    1721                 "The number of FixedLocator locations"
    1722                 f" ({len(locator.locs)}), usually from a call to"

```

ValueError: The number of FixedLocator locations (3), usually from a call to set_ticks, does not match the number of ticklabels (1350).



```
In [94]: df.rename(columns={'Covid- infected or not':'Covid- infected'}, inplace=True)
```

```
In [95]: #To Visualize data
import seaborn as sns
#To partition the data
from sklearn.model_selection import train_test_split
#Importing library for logistic regression
from sklearn.linear_model import LogisticRegression
#Importing performance metrics -accuracy score & confusion matrix
from sklearn.metrics import accuracy_score, confusion_matrix

from sklearn.ensemble import ExtraTreesClassifier
```

```
In [96]: feature_cols = ['Age', 'CAS2', 'CAS3', 'CAS4', 'CAS5', 'Res. No', 'Education', 'Income', 'Gender', 'Marital status']
X = df[feature_cols]
y = df['Covid- infected'].astype('float')

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)


#create a dict for final graph
#Use : methodDict['Stacking'] = accuracy_score
methodDict = {}
rmseDict = ()
```

```
In [97]: forest = ExtraTreesClassifier(n_estimators=250, random_state=0)

forest.fit(X,y)
importances = forest.feature_importances_
std = np.std([tree.feature_importances_ for tree in forest.estimators_],
             axis=0)
indices = np.argsort(importances)[::-1]

labels = []
for f in range(X.shape[1]):
    labels.append(feature_cols[f])

plt.figure(figsize=(12,8))
plt.title("Feature importances")
plt.bar(range(X.shape[1]), importances[indices],
        color = 'r', yerr=std[indices], align="center")
plt.xticks(range(X.shape[1]), labels, rotation='vertical')
plt.xlim([-1,X.shape[1]])
plt.show()
```



Evaluating models

```
In [98]: """Logistic Regression"""
```

```
Out[98]: 'Logistic Regression'
```

```
In [99]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, mean_squared_error, precision_recall_curve
from sklearn.model_selection import cross_val_score
from sklearn import metrics

from scipy.stats import randint

#Importing performance metrics -accuracy score & confussion matrix
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
In [100]: data = df.copy()
```

```
In [101]: data.shape
```

```
Out[101]: (1350, 12)
```

```
In [102]: data["Covid- infected"] = pd.get_dummies(data["Covid- infected"])
```

```
-----  
ValueError                                Traceback (most recent call last)  
~\AppData\Local\Temp\ipykernel_8180\934390265.py in <module>  
----> 1 data["Covid- infected"] = pd.get_dummies(data["Covid- infected"])  
  
~\anaconda3\lib\site-packages\pandas\core\frame.py in __setitem__(self, key, value)  
    3643         self._setitem_array(key, value)  
    3644     elif isinstance(value, DataFrame):  
-> 3645         self._set_item_frame_value(key, value)  
    3646     elif (  
    3647         is_list_like(value)  
  
~\anaconda3\lib\site-packages\pandas\core\frame.py in _set_item_frame_value(self, key, value)  
    3773         len_cols = 1 if is_scalar(cols) else len(cols)  
    3774         if len_cols != len(value.columns):  
-> 3775             raise ValueError("Columns must be same length as key")  
    3776  
    3777         # align right-hand-side columns if self.columns
```

ValueError: Columns must be same length as key

```
In [103]: data.shape
```

```
Out[103]: (1350, 12)
```

```
In [104]: new_data = pd.get_dummies(data , drop_first = True)
new_data.head()
```

Out[104]:

	Res. No	Age	Gender	Marital status	Education	Income	Covid- infected	CAS2	CAS3	CAS4	CAS5	CAS1_1	CAS1_2	CAS1_3	CAS1_c
0	1	2	1	1	1	1	1	1	2	2	2	1	0	0	0
1	2	1	0	1	1	2	1	1	2	2	2	1	0	0	0
2	3	2	0	1	3	1	1	3	0	0	3	0	0	1	0
3	4	3	0	0	1	2	1	1	0	1	0	0	0	0	0
4	5	2	0	1	1	1	1	2	1	1	2	1	0	0	0

```
In [105]: columns_list = list(new_data.columns)
print(columns_list)
```

```
['Res. No', 'Age', 'Gender', 'Marital status', 'Education', 'Income', 'Covid- infected', 'CAS2', 'CAS3', 'CAS4', 'CAS5', 'CAS1_1', 'CAS1_2', 'CAS1_3', 'CAS1_c']
```

```
In [106]: features = list(set(columns_list)-set(['Covid- infected']))
print(features)
```

```
['Income', 'CAS1_3', 'CAS5', 'CAS1_1', 'CAS1_2', 'CAS3', 'CAS1_c', 'Res. No', 'Age', 'Gender', 'CAS2', 'CAS4', 'Education', 'Marital status']
```

```
In [107]: y = new_data['Covid- infected'].values
print(y)
```

```
[1 1 1 ... 0 1 0]
```

```
In [108]: x = new_data[features].values  
print(x)
```

```
[[1 0 2 ... 2 1 1]  
 [2 0 2 ... 2 1 1]  
 [1 1 3 ... 0 3 1]  
 ...  
 [2 0 2 ... 1 1 0]  
 [2 1 3 ... 2 1 0]  
 [1 0 1 ... 0 1 1]]
```

```
In [109]: x.shape
```

```
Out[109]: (1350, 14)
```

```
In [110]: y.shape
```

```
Out[110]: (1350,)
```

```
In [111]: train_x, test_x, train_y, test_y = train_test_split(x, y, test_size=0.3, random_state=0)
```

```
In [112]: train_x.shape
```

```
Out[112]: (945, 14)
```

```
In [113]: train_y.shape
```

```
Out[113]: (945,)
```

```
In [114]: test_x.shape
```

```
Out[114]: (405, 14)
```

```
In [115]: test_y.shape
```

```
Out[115]: (405,)
```

```
In [116]: #make instance of the model  
logistic = LogisticRegression()
```

```
In [117]: logistic.fit(train_x,train_y)  
logistic.coef_
```

C:\Users\balakumar\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning:

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
Out[117]: array([[ 4.22222560e-01,  1.12845534e-01, -1.00519405e-01,  
                  1.74817364e-01, -1.19710481e-01,  2.91020158e-03,  
                  1.42915016e-02,  8.83981370e-05, -1.06753598e-02,  
                 -4.02282521e-03,  5.81813176e-02, -7.39898001e-02,  
                  1.45447564e-01,  7.52157425e-01]])
```

```
In [45]: logistic.intercept_
```

```
Out[45]: array([0.41789231])
```



```
In [50]: from sklearn.ensemble import RandomForestClassifier
```

```
In [51]: classifierclf = RandomForestClassifier()
```

```
In [52]: classifierclf.fit(x_train,y_train)
```

```
Out[52]: RandomForestClassifier()
```

```
In [53]: x_train.shape
```

```
Out[53]: (1080, 14)
```

```
In [54]: y_pred = classifierclf.predict(x_test)
```

```
In [55]: #perform standziation  
from sklearn.metrics import accuracy_score  
accuracy_score(y_test,y_pred)
```

```
Out[55]: 0.8296296296296296
```

PCA

```
In [56]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()
```

```
In [57]: x_train = scaler.fit_transform(x_train)  
x_test = scaler.transform(x_test)
```

```
In [58]: from sklearn.decomposition import PCA  
pca = PCA(n_components = 10)
```

```
In [59]: x_train_trf = pca.fit_transform(x_train)
x_test_trf = pca.transform(x_test)
```

```
In [60]: x_train_trf.shape
```

```
Out[60]: (1080, 10)
```

```
In [61]: x_train_trf.shape
```

```
Out[61]: (1080, 10)
```

```
In [62]: classifierclf = RandomForestClassifier()
```

```
In [63]: classifierclf.fit(x_train_trf,y_train)
```

```
Out[63]: RandomForestClassifier()
```

```
In [64]: y_pred = classifierclf.predict(x_test_trf)
```

```
In [65]: accuracy_score(y_test,y_pred)
```

```
Out[65]: 0.8185185185185185
```


```
In [66]: pca = PCA(n_components=3)
x_train_trf = pca.fit_transform(x_train)
x_test_trf = pca.transform(x_test)
```



```
In [67]: x_train_trf
```

```
Out[67]: array([[ -0.800831   ,  0.46627249,  1.81862815],  
                [ -1.45561388, -0.14720139,  0.83843941],  
                [ -0.06217479,  0.66876324,  1.84793879],  
                ...,  
                [  4.34139027,  0.2973388 ,  0.25442508],  
                [ -0.4374745 ,  0.94716784,  0.58549119],  
                [ -2.27695521,  0.27376842,  0.75364073]])
```

```
In [68]: import plotly.express as px
y_train_trf = y_train.astype(str)
fig = px.scatter(x=x_train_trf[:,0],
                 y=x_train_trf[:,1],
                 color = y_train_trf,
                 color_discrete_sequence = px.colors.qualitative.G10
                )
fig.show()
```




```
In [69]: pca = PCA(n_components=3)
x_train_trf = pca.fit_transform(x_train)
x_test_trf = pca.transform(x_test)
```

```
In [70]: x_train_trf
```

```
Out[70]: array([[ -0.80036973,  0.46743209,  1.82025925],  
                [ -1.45580267, -0.14723885,  0.83708358],  
                [ -0.06187977,  0.66949332,  1.84898979],  
                ...,  
                [  4.34148159,  0.29751105,  0.25482023],  
                [ -0.43788173,  0.94660988,  0.58329507],  
                [ -2.2771019 ,  0.27388197,  0.75239641]])
```

```
In [71]: y_train_trf = y_train.astype(str)
fig = px.scatter_3d(data, x=x_train_trf[:,0],
                    y=x_train_trf[:,1],
                    z=x_train_trf[:,2],
                    color = y_train_trf)

fig.update_layout(margin=dict(l=20, r=20, t=20, b=20))
fig.show()
```



```
In [72]: #Eigen values  
pca.explained_variance_
```

```
Out[72]: array([1.87761549, 1.6960011 , 1.37328951])
```

```
In [73]: #finding optimum number of principle components  
#to check percentage of variance explained by individual eigen vector  
  
pca.explained_variance_ratio_  
  
#eigen vector 1:13%  
#eigen vector 2:12%  
#eigen vector 3:09%.....
```

```
Out[73]: array([0.13399121, 0.12103077, 0.09800128])
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```